

Classification Systems

Daniel Cerdán, Fernando Freire

March 27, 2019

Contents

1	Classification Systems	2
1.1	Methods	3
1.1.1	Implementation details	3
1.2	Breast cancer	13
1.2.1	Hyper-parameter guessing plots	16
1.3	Prostate cancer	20
1.3.1	Hyper-parameter guessing plots	24
1.4	Conclusions	28

1 Classification Systems

In this practical, you are asked to compare the prediction error of:

1. The Naive Bayes Classifier
2. LDA
3. QDA
4. Nearest Shrunk Centroids Classifier

On the Breast Cancer dataset provided in the previous notebooks, and the Prostate cancer dataset attached. The details about this last dataset are found in the reference:

Singh, D., Febbo, P., Ross, K., Jackson, D., Manola, J., Ladd, C., Tamayo, P., Renshaw, A., D'Amico, A., Richie, J., Lander, E., Loda, M., Kantoff, P., Golub, T., & Sellers, W. (2002). Gene expression correlates of clinical prostate cancer behavior. *Cancer Cell*, 1, 203–209.

This dataset is in CSV format and the last column contains the class label. The task of interest is to discriminate between normal and tumor tissue samples.

Importantly:

Use a random split of 2 / 3 of the data for training and 1 / 3 for testing each classifier. Any hyper-parameter of each method should be tuned using a grid-search guided by an inner cross-validation procedure that uses only training data. To reduce the variance of the estimates, report average error results over 20 different partitions of the data into training and testing as described above. Submit a notebook showing the code and the results obtained. Give some comments about the results and respond to these questions:

What method performs best on each dataset? What method is more flexible? What method is more robust to over-fitting?

Script 1.0.1 (python)

```
1 import warnings
2 warnings.filterwarnings("ignore")
3 %matplotlib inline
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import matplotlib.lines as mlines
8 import matplotlib as mpl
9 from matplotlib import colors
10 import seaborn as sns; sns.set()
11 import scipy.stats as stats
12 import scipy as sp
13 from scipy import linalg
14 from sklearn.naive_bayes import GaussianNB
15 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
16 from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
17 from sklearn.naive_bayes import GaussianNB
18 from sklearn.neighbors import NearestCentroid
19 from sklearn.decomposition import PCA
20 from sklearn.pipeline import Pipeline
21 from sklearn.model_selection import train_test_split, RepeatedStratifiedKFold, GridSearchCV
22 from sklearn import preprocessing
23 from sklearn.metrics import accuracy_score, make_scorer, confusion_matrix,
    → classification_report, precision_score
```

1.1 Methods

These are the python methods that encapsulate the four learning methods.

1.1.1 Implementation details

Quadratic Discriminant Analysis

Before training the classifier we have chosen a good value for the corresponding regularization hyper-parameter with a grid-search guided by cross-validation.

The regularization parameter regularizes the covariance matrix estimate as

$$(1 - \lambda) \cdot \Sigma + \lambda \cdot \mathbf{I}$$

Nearest Centroids

Before training the classifier we have chosen a good value for the shrinkage threshold hyper-parameter with a grid-search guided by cross-validation.

This procedure leads to a reduction in the number of features, by zeroing all deltas that exceed the threshold.

They take the form:

$$\mu_{kj} = m_j + \Delta_{kj},$$

where Δ_{kj} is the shrunken component

Selecting the best parameter value

To do so, as a first idea, we compute the set of values with the maximum test data accuracy, and between then we choose the set of values that have the maximum train data accuracy. From this set we choose the lowest value.

We compare this method defined in the global parameter `gp_best_hyper_method = "max_in_test"` to the method that simply computes the parameter for maximum cross validation data accuracy (`gp_best_hyper_method = "max_in_cv"`) and we obtain better accuracies in test data. But we came into account that because we use test data to adjust the hyper-parameter perhaps we are introducing some bias improving the accuracy of this particular test data, and so a optimistic evaluation of the accuracy.

Thus, we compute the hyperparameters with "max_in_cv" choose method.

Script 1.1.1 (python)

```
1 # Global parameters
2 # Verbose flag
3 gp_verbose = False
4 # Show progress flag
5 gp_show_progress = True
6 # Disable plots
7 gp_disable_plots = True
8 # Activate QDA with hyper-parameter reg_param
9 gp_qda_hyper = True
10 # Activate NSC with hyper-parameter shrink_threshold
11 gp_nsc_hyper = True
12 # Dimensionality reduction(PCA)
13 gp_dim_red = False
14 # Retained variance (PCA)
15 gp_retained_variance = 99
16 # Number of iterations
17 gp_iterations = 1
```

```

18 # Test size = number of samples / gp_test_size
19 gp_test_size = 3
20 # Best_hyper_param_method
21 gp_best_hyper_method = "max_in_cv"
22 # Skfold splits for hyperparameter guessing
23 gp_skfold_splits = 10
24
25 # Global execution parameters
26 # New dimensions after PCA
27 ge_features_reduction = 0
28
29 #Number of params of learning methods
30 ge_complexity = []
31
32 # Methods
33 def get_component_number(df_data, desired_variance=99.0, scaling=False):
34     """
35     Obtain the number of components that explains a %desired_variance
36     Args:
37         df_data (dataframe): dataframe of features in cols and samples in rows
38         desired_variance (float): desired explained variance
39         scaling (boolean): True if pre-scaling is needed prior to compute PCA
40     Returns:
41         int: number of components to maintain to have a explained variance >=
→ desired_variance
42         float: variance explained for the number of components returned
43         numpy array: cumulative variance by number of components retained
44     """
45     if scaling:
46         df_data_2 = preprocessing.StandardScaler().fit_transform(df_data)
47     else:
48         df_data_2 = df_data
49     # project the data into this new PCA space
50     pca = PCA().fit(df_data_2)
51     desired_variance = desired_variance/100.0
52     explained_variance = np.cumsum(pca.explained_variance_ratio_)
53     component_number = 0
54     for cumulative_variance in explained_variance:
55         component_number += 1
56         if cumulative_variance >= desired_variance:
57             break
58     return component_number, cumulative_variance, explained_variance
59
60
61 def create_datasets_from_file(data_file, header, random_state, label_pos,
62                             label_value, features_ini, features_fin=None,
63                             with_dim_red=False, retained_variance=99.0,
64                             reuse=False, dataset=None, labels=None):
65     """Create training and test sets from file
66
67     Args:
68         data_file (string): Name of the data file (csv) of samples a features

```

```

69         header (string): None or position of the header (pandas read_csv parameter)
70         random_state (int): Seed for the random split (as needed for sklearn
→ train_test_split)
71         label_pos (int): Column of the labels in data_file
72         label_value (int): Value of the label to assign internal '1' value
73         features_ini (int): First column of features in data_file
74         features_fin (int): Last column + 1 of features in data_file. If None, last
→ column of file.
75         with_dim_red (bool): If True, it performs a dimensionality reduction by PCA
76         retained_variance (float): If dimensionality reduction, variance to retain
77         reuse (bool): Reuse previous dataset
78         dataset: Dataset to reuse
79         labels: Labels to reuse
80
81     Returns:
82         (np.array): train set scaled
83         (np.array): test set scaled
84         (np.array): class labels for the train set
85         (np.array): class labels for the test set
86         (np.array): dataset
87         (np.array): labels
88
89     """
90     global ge_features_reduction
91     if not reuse:
92         data = pd.read_csv(data_file, header = header)
93         if features_fin == None:
94             X = data.values[ :, features_ini:].astype(np.float)
95         else:
96             X = data.values[ :, features_ini:features_fin].astype(np.float)
97             y = (data.values[ :, label_pos ] == label_value).astype(np.int)
98     else: #reuse previous dataset
99         X = dataset
100         y = labels
101
102     # Split dataset between training and test
103     x_train, x_test, y_train, y_test = train_test_split(X, y,
104                                                         test_size=1.0/gp_test_size,
→ random_state=random_state)
105
106     # Data standardization
107     scaler = preprocessing.StandardScaler().fit(x_train)
108     x_train_scaled = scaler.transform(x_train)
109     x_test_scaled = scaler.transform(x_test)
110     # Check standardization
111     for i in range (1, np.size(x_train_scaled,1)):
112         assert round(np.var(x_train_scaled[:,0]),3) == round(np.var(x_train_scaled[:,i]),3),\
113             "Warning: revise data standardization"
114
115     if with_dim_red:
116         desired_variance = retained_variance
117         component_number, _, _ =\
            get_component_number(x_train_scaled, desired_variance, scaling=None)

```

```

118         if gp_verbose: print("Features reduced to", component_number)
119         ge_features_reduction = component_number
120         pca = PCA(n_components = component_number)
121         pca.fit(x_train_scaled)
122         x_train_scaled = pca.transform(x_train_scaled)
123         x_test_scaled = pca.transform(x_test_scaled)
124
125     return x_train_scaled, x_test_scaled, y_train, y_test, X, y
126
127 def prediction_accuracy(x_train, x_test, y_train, y_test, method_func, method_param="",
→ param_value=""):
128     """Estimate parameter given training and test sets:
129     Args:
130         x_train (np.array): train set
131         x_test (np.array): test set
132         y_train (np.array): class labels for the train set
133         y_test (np.array): class labels for the test set
134         method_func (string) : name of the learning method
135         method_param (string): name of learning method parameter
136         param_value (float): value of parameter to try
137     Returns:
138         float: best parameter value to use in prediction
139
140     """
141     if method_param != "" :
142         params = {method_param : param_value}
143     else:
144         params = {}
145     method = globals()[method_func](**params)
146
147     # Training
148     method.fit(x_train, y_train)
149
150     ge_complexity.append([method_func, len(method.get_params())])
151     # Prediction of test
152     y_pred = method.predict(x_test)
153     conf_test = confusion_matrix(y_test, y_pred, labels=[1,0])
154
155     # Prediction of train
156     y_pred = method.predict(x_train)
157     conf_train = confusion_matrix(y_train, y_pred, labels=[1,0])
158
159     return conf_train, conf_test
160
161 def estimate_parameter(x_train, x_test, y_train, y_test,
162                       method_func, param, param_values,
163                       best_param_value_method="max_in_test"):
164     """Estimate parameter given training and test sets:
165     Args:
166         x_train (np.array): train set
167         x_test (np.array): test set
168         y_train (np.array): class labels for the train set

```

```

169         y_test (np.array): class labels for the test set
170         method_func (string) : name of the learning method
171         param (string): name of learning method parameter
172         param_values (list of float): list of parameter values to try
173         best_param_value_method: if "max_in_test" gives the value with the maximum
→ accuracy
174                                     in test data.
175
176     Returns:
177         (float): best parameter value to use in prediction
178
179     """
180     # Pipeline for estimate the regularization parameter
181     pipeline = Pipeline([ ('method', globals()[method_func]() )])
182
183     # Construct the grid the hyperparameter candidate shronk theshold
184     param_grid = { 'method__' + param : param_values }
185
186     # Evaluating
187     skfold = RepeatedStratifiedKfold(n_splits=gp_skfold_splits, n_repeats=1, random_state=0)
188     gridcv = GridSearchCV(pipeline, cv=skfold, n_jobs=1, param_grid=param_grid,\
189                           scoring=make_scorer(accuracy_score))
190     result = gridcv.fit(x_train, y_train)
191
192     # Accuracies
193     accuracies = gridcv.cv_results_['mean_test_score']
194     std_accuracies = gridcv.cv_results_['std_test_score']
195
196     test_accuracies = np.ones(len(param_values))
197
198     for i in range(len(param_values)):
199         method_params = {param : param_values[ i ]}
200         method = globals()[method_func](**method_params)
201         method.fit(x_train, y_train)
202         test_accuracies[ i ] = accuracy_score(method.predict(x_test), y_test)
203
204     # Obtain best_param_value as max
205     max_test_accuracy = max(test_accuracies)
206     if best_param_value_method == "max_in_test":
207         best_param_value = 0
208         best_train_accuracy = 0
209         for i in range(len(param_values)):
210             if test_accuracies[ i ] == max_test_accuracy:
211                 if accuracies[i] > best_train_accuracy:
212                     best_train_accuracy = accuracies[i]
213                     best_param_value = param_values[i]
214     else:
215         best_param_value = param_values[ np.argmax(accuracies) ]
216
217     # Plot
218     if not gp_disable_plots:
219         plt.figure(figsize=(9, 9))
220         line1, = plt.plot(param_values, accuracies, 'o-', color="g")
221         line2, = plt.plot(param_values, test_accuracies, 'x-', color="r")

```

```

220 plt.fill_between(param_values, accuracies - std accuracies / np.sqrt(10), \
221                  accuracies + std accuracies / np.sqrt(10), alpha=0.1, color="g")
222 plt.grid()
223 plt.title("Different hyper-parameter " + param + " values for " + method_func)
224 plt.xlabel('Hyper-parameter')
225 plt.xticks(np.round(np.array(param_values), 2))
226 plt.ylabel('Classification Accuracy')
227 plt.ylim((min(min(accuracies), min(test accuracies)) - 0.1,
228           min(1.02, max(max(accuracies), max(test accuracies)) + 0.1)))
229
230 plt.xlim((min(param_values), max(param_values)))
231 legend_handles = [ mlines.Line2D([], [], color='g', marker='o', \
232                             markersize=15, label='CV-estimate'), \
233                   mlines.Line2D([], [], color='r', marker='x', \
234                             markersize=15, label='Test set estimate')]
235 plt.legend(handles=legend_handles, loc = 3)
236 plt.show()
237 if gp_verbose:
238     print("Best param value %s Method %s: %s" % (method_func, best_param_value_method,
239         → best_param_value))
239 return best_param_value
240
241 def calculate_avg_results(train_results, test_results, classifier_name):
242     """Calculate the average accuracy, TPN and TNR for the n_gp_iterations of a classifier
243     Args:
244         train_results (array nx2x2): Each of the n confusions matrix generated for the
245     → train set
246         test_results (array nx2x2): Each of the n confusions matrix generated for the
247     → test set
248         classifier_name (string): Classifier name gp_iterations print the results
249     Returns:
250         (np.array 3): Average accuracy, TPR and TNR of the n iterations of the train set
251         (np.array 3): Average accuracy, TPR and TNR of the n iterations of the test set
252         (np.array 3): Stdev accuracy, TPR and TNR of the n iterations of the train set
253         (np.array 3): Stdev accuracy, TPR and TNR of the n iterations of the test set
254
255     """
256
257     measures_train = np.zeros(shape = (gp_iterations,3)) # Each row is an ex. and each
258     → column is the accuracy, TPR and TNR
259     measures_test = np.zeros(shape =(gp_iterations,3))
260
261     avg_results_train = np.zeros(3) # Each element is the average accuracy, the TPR and the
262     → TNR
263     avg_results_test = np.zeros(3)
264     # TP in 0,0
265     # FN in 0,1
266     # TN in 1,1
267     # FP in 1,0
268     for i in range(gp_iterations):
269         # For train set
270         TN = train_results[i,1,1]

```



```

267     TP = train_results[i,0,0]
268     FP = train_results[i,1,0]
269     FN = train_results[i,0,1]
270     measures_train[i][0] = (TP + TN) / (TN + TP + FP + FN)
271     measures_train[i][1] = (TP / (TP + FN))
272     measures_train[i][2] = (TN / (TN + FP))
273
274     #For the test set
275     TN = test_results[i,1,1]
276     TP = test_results[i,0,0]
277     FP = test_results[i,1,0]
278     FN = test_results[i,0,1]
279     measures_test[i][0] = (TP + TN) / (TN + TP + FP + FN)
280     measures_test[i][1] = (TP / (TP + FN))
281     measures_test[i][2] = (TN / (TN + FP))
282
283     avg_results_train = measures_train.mean(axis = 0)
284     avg_results_test = measures_test.mean(axis = 0)
285     std_results_train = measures_train.std(axis = 0)
286     std_results_test = measures_test.std(axis = 0)
287     if gp_verbose:
288         print("\nResults for the ", classifier_name, " classifier\n")
289
290         print('For the train set:')
291         print('Prediction accuracy of train set is: %f' % avg_results_train[0])
292         print('True positive rate of train set is: %f' % avg_results_train[1])
293         print('True negative rate of train set is: %f\n' % avg_results_train[2])
294
295         print('For the test set:')
296         print('Prediction accuracy of train set is: %f' % avg_results_test[0])
297         print('True positive rate of train set is: %f' % avg_results_test[1])
298         print('True negative rate of train set is: %f\n' % avg_results_test[2])
299
300     return avg_results_train, avg_results_test, std_results_train, std_results_test
301
302 def print_parameters():
303     print("Parameters")
304     pdata = {'Parameters' : [gp_qda_hyper, gp_nsc_hyper,
305                             gp_dim_red, gp_retained_variance,
306                             gp_iterations, 100.0/gp_test_size,
307                             gp_best_hyper_method]}
308     pdataf = pd.DataFrame(data = pdata, index = ['QDA hyper-parameter reg_param',
309                                                'NSC Hyper-parameter shrink_threshold',
310                                                'Dimensionality Reduction',
311                                                'Retained variance',
312                                                'Number of iterations',
313                                                'Test set size',
314                                                'Best hyperparameter select method'])
315     display(pdataf)
316     print("")
317
318 def print_execution_data():

```

```

319 print("Execution data")
320 data = [ge_samples, ge_features, ge_features_reduction]
321 index = ['Number of samples', 'Number of features', 'Features reduction']
322 for [method, complexity] in ge_complexity:
323     data.append(complexity)
324     index.append("Params " + method)
325 pdata = {'Execution Data' : data}
326 pdataf = pd.DataFrame(data = pdata, index = index)
327 display(pdataf)
328 print("")
329
330 def print_accuracies(accuracy_NBC, accuracy_LDA, accuracy_QDA, accuracy_NSC, metric):
331     print(metric)
332     d = {}
333     for m in [accuracy_NBC, accuracy_LDA, accuracy_QDA, accuracy_NSC]:
334         overfit_indicator = 1000 * round(abs(m[1] - m[0]) / (m[0] + 0.0000001), 4)
335         #overfit_indicator = round(abs(m[1] - m[0]),3)
336         m[0] = (round(m[0],2), round(m[2],2))
337         m[1] = (round(m[1],2), round(m[3],2))
338         m[2] = overfit_indicator
339     d = {'NBC': accuracy_NBC[:3],
340         'LDA': accuracy_LDA[:3],
341         'QDA': accuracy_QDA[:3],
342         'NSC': accuracy_NSC[:3]}
343     df = pd.DataFrame(data = d, index = ['Train', 'Test', 'Overfit degree'])
344     display(df)
345     print("")
346
347 def learn_dataset(data_file, header, random_state, label_pos,
348                 label_value, features_ini, features_fin=None,
349                 best_param_value_method="max_in_test",
350                 with_dim_red=False, retained_variance=99.0):
351     """Learn data sets from file, methods:
352     1. The Naive Bayes Classifier
353     2. LDA
354     3. QDA
355     4. Nearest Shrunken Centroids Classifier
356     Args:
357     data_file (string): Name of the data file (csv) of samples a features
358     header (string): None or position of the header (pandas read_csv parameter)
359     random_state (int): Seed for the random split of sets (as needed for sklearn
360     train_test_split)
361     label_pos (int): Column of the labels in data_file
362     label_value (int): Value of the label to assign internal '1' value. We consider
363     this label as
364     the positive label in prediction validation. We assign malign or cancer status to
365     this label.
366     features_ini (int): First column of features in data_file
367     features_fin (int): Last column + 1 of features in data_file. If None, last
368     column of file
369     best_param_value_method (str): if "max_in_test" gives the value with the maximum
370     accuracy

```

```

366                                     in test data
367         with_dim_red (bool): If True, it performs a dimensionality reduction by PCA
368         retained_variance (float): If dimensionality reduction, variance to retain
369
370     """
371     global ge_features_reduction
372     global ge_features
373     global ge_samples
374     global ge_complexity
375     ge_features_reduction = 0
376     nbc_train = np.zeros(shape=(gp_iterations,2,2))
377     nbc_test = np.zeros(shape=(gp_iterations,2,2))
378     lda_train = np.zeros(shape=(gp_iterations,2,2))
379     lda_test = np.zeros(shape=(gp_iterations,2,2))
380     qda_train = np.zeros(shape=(gp_iterations,2,2))
381     qda_test = np.zeros(shape=(gp_iterations,2,2))
382     nsc_train = np.zeros(shape=(gp_iterations,2,2))
383     nsc_test = np.zeros(shape=(gp_iterations,2,2))
384
385
386     for i in range(gp_iterations):
387         ge_complexity = []
388         if gp_show_progress: print("\nIteration: ",i)
389         if i == 0:
390             reuse = False
391             X = None
392             y = None
393         else:
394             reuse = True
395         X_train_scaled, X_test_scaled, y_train, y_test, X, y = \
396             create_datasets_from_file(data_file, header, random_state + i,
397                                     label_pos, label_value, features_ini, features_fin =
398                                     → features_fin,
399                                     with_dim_red = with_dim_red, retained_variance =
400                                     → retained_variance,
401                                     reuse = reuse, dataset = X, labels = y)
402
403         ge_samples = X.shape[0]
404         ge_features = X.shape[1]
405         if gp_verbose: print(X_train_scaled.shape)
406
407         if gp_verbose: print("NBC")
408         # Naive Bayes accuracy
409         nbc_train[i], nbc_test[i] = prediction_accuracy(X_train_scaled, X_test_scaled,
410                                                         → y_train, y_test, "GaussianNB")
411
412         # LDA accuracy
413         if gp_verbose: print("LDA")
414         lda_train[i], lda_test[i] = prediction_accuracy(X_train_scaled, X_test_scaled,
415                                                         → y_train, y_test, "LinearDiscriminantAnalysis")
416
417         # QDA estimate reg parameter
418         if gp_verbose: print("QDA")

```

```

414 if gp_qda_hyper:
415     param_values = np.linspace(0, 1, 10).tolist()
416     best_param_value = estimate_parameter(X_train_scaled, X_test_scaled, y_train,
417     ↪ y_test,\
418     ↪ "QuadraticDiscriminantAnalysis", "reg_param", param_values,\
419     ↪ best_param_value_method)
420     # QDA accuracy
421     # Best parameter reg value according CV estimate
422     qda_train[i], qda_test[i] = prediction_accuracy(X_train_scaled, X_test_scaled,
423     ↪ y_train, y_test,\
424     ↪ "QuadraticDiscriminantAnalysis", "reg_param",
425     ↪ best_param_value)
426 else:
427     qda_train[i], qda_test[i] = prediction_accuracy(X_train_scaled, X_test_scaled,
428     ↪ y_train, y_test,\
429     ↪ "QuadraticDiscriminantAnalysis")
430 # Centroids
431 if gp_verbose: print("NSC")
432 if gp_nsc_hyper:
433     # Best parameter shrink_threshold value according CV estimate
434     param_values = np.linspace(0, 8, 20).tolist()
435     best_param_value = estimate_parameter(X_train_scaled, X_test_scaled, y_train,
436     ↪ y_test,\
437     ↪ "NearestCentroid", "shrink_threshold", param_values,\
438     ↪ best_param_value_method)
439     # Centroids accuracy
440     nsc_train[i], nsc_test[i] = prediction_accuracy(X_train_scaled, X_test_scaled,
441     ↪ y_train, y_test,\
442     ↪ "NearestCentroid", "shrink_threshold",
443     ↪ best_param_value)
444 else:
445     nsc_train[i], nsc_test[i] = prediction_accuracy(X_train_scaled, X_test_scaled,
446     ↪ y_train, y_test,\
447     ↪ "NearestCentroid")
448 # Calculate and print the average results
449 avg_results_train_NBC, avg_results_test_NBC, std_results_train_NBC, std_results_test_NBC
450 ↪ = calculate_avg_results(nbc_train, nbc_test, "NBC")
451 avg_results_train_LDA, avg_results_test_LDA, std_results_train_LDA, std_results_test_LDA
452 ↪ = calculate_avg_results(lda_train, lda_test, "LDA")
453 avg_results_train_QDA, avg_results_test_QDA, std_results_train_QDA, std_results_test_QDA
454 ↪ = calculate_avg_results(qda_train, qda_test, "QDA")
455 avg_results_train_NSC, avg_results_test_NSC, std_results_train_NSC, std_results_test_NSC
456 ↪ = calculate_avg_results(nsc_train, nsc_test, "NSC")
457
458 print_parameters()
459 i = 0
460 for metric in ["Accuracy", "TPR", "TNR"]:
461     print_accuracies([avg_results_train_NBC[i], avg_results_test_NBC[i],
462     ↪ std_results_train_NBC[i], std_results_test_NBC[i]],
463     ↪ [avg_results_train_LDA[i], avg_results_test_LDA[i],
464     ↪ std_results_train_LDA[i], std_results_test_LDA[i]],

```

```

451         [avg_results_train_QDA[i], avg_results_test_QDA[i],
452          ↪ std_results_train_QDA[i], std_results_test_QDA[i]],
453         [avg_results_train_NSC[i], avg_results_test_NSC[i],
454          ↪ std_results_train_NSC[i], std_results_test_NSC[i]], metric)
453     i += 1
454     print_execution_data()

```

1.2 Breast cancer

Script 1.2.1 (python)

```

1  gp_disable_plots = True
2  gp_verbose = False
3  gp_dim_red = False
4  gp_retained_variance = 99
5  gp_qda_hyper = True
6  gp_nsc_hyper = True
7  gp_iterations = 20
8  gp_show_progress = False
9  gp_best_hyper_method = "max_in_cv"
10 gp_skfold_splits = 10
11
12 learn_dataset(data_file = './data/wdbc.csv', header = None, random_state=0,
13               label_pos=1, label_value="M", features_ini = 2, features_fin = None,
14               best_param_value_method = gp_best_hyper_method,
15               with_dim_red = gp_dim_red, retained_variance = gp_retained_variance)
16
17 gp_dim_red = True
18 learn_dataset(data_file = './data/wdbc.csv', header = None, random_state=0,
19               label_pos=1, label_value="M", features_ini = 2, features_fin = None,
20               best_param_value_method = gp_best_hyper_method,
21               with_dim_red = gp_dim_red, retained_variance = gp_retained_variance)

```

Output

Parameters

	Parameters
QDA hyper-parameter reg_param	True
NSC Hyper-parameter shrink_threshold	True
Dimensionality Reduction	False
Retained variance	99
Number of iterations	20
Test set size	33.3333
Best hyperparameter select method	max_in_cv

Output

Accuracy

	NBC	LDA	QDA	NSC
Train	(0.94, 0.01)	(0.97, 0.01)	(0.97, 0.01)	(0.94, 0.01)
Test	(0.93, 0.01)	(0.96, 0.01)	(0.97, 0.01)	(0.94, 0.01)
Overfit degree	5.4	7.5	9.4	7

Output

TPR

	NBC	LDA	QDA	NSC
Train	(0.9, 0.01)	(0.92, 0.01)	(0.95, 0.01)	(0.9, 0.02)
Test	(0.9, 0.03)	(0.89, 0.03)	(0.93, 0.03)	(0.9, 0.04)
Overfit degree	10.1	23.3	14.7	3.6

Output

TNR

	NBC	LDA	QDA	NSC
Train	(0.97, 0.01)	(1.0, 0.0)	(0.99, 0.0)	(0.97, 0.01)
Test	(0.95, 0.02)	(1.0, 0.0)	(0.98, 0.02)	(0.96, 0.02)
Overfit degree	14.3	0	6.8	9.2

Output

Execution data

	Execution Data
Number of samples	569
Number of features	30
Features reduction	0
Params GaussianNB	2
Params LinearDiscriminantAnalysis	6
Params QuadraticDiscriminantAnalysis	5
Params NearestCentroid	2

Output

Parameters

	Parameters
QDA hyper-parameter reg_param	True
NSC Hyper-parameter shrink_threshold	True
Dimensionality Reduction	True
Retained variance	99
Number of iterations	20
Test set size	33.3333
Best hyperparameter select method	max_in_cv

Output

Accuracy

	NBC	LDA	QDA	NSC
Train	(0.9, 0.01)	(0.96, 0.0)	(0.98, 0.01)	(0.93, 0.01)
Test	(0.89, 0.02)	(0.96, 0.01)	(0.97, 0.01)	(0.93, 0.02)
Overfit degree	10.2	7.4	8.8	5.7

Output

TPR

	NBC	LDA	QDA	NSC
Train	(0.84, 0.02)	(0.91, 0.01)	(0.95, 0.02)	(0.87, 0.01)
Test	(0.83, 0.04)	(0.89, 0.02)	(0.94, 0.03)	(0.87, 0.03)
Overfit degree	5.4	18.6	17.2	4

Output

TNR

	NBC	LDA	QDA	NSC
Train	(0.94, 0.01)	(1.0, 0.0)	(0.99, 0.0)	(0.97, 0.01)
Test	(0.93, 0.03)	(0.99, 0.01)	(0.99, 0.01)	(0.96, 0.02)
Overfit degree	13.1	2.1	4.6	7.6

Output

Execution data


	Execution Data
Number of samples	569
Number of features	30
Features reduction	17
Params GaussianNB	2
Params LinearDiscriminantAnalysis	6
Params QuadraticDiscriminantAnalysis	5
Params NearestCentroid	2

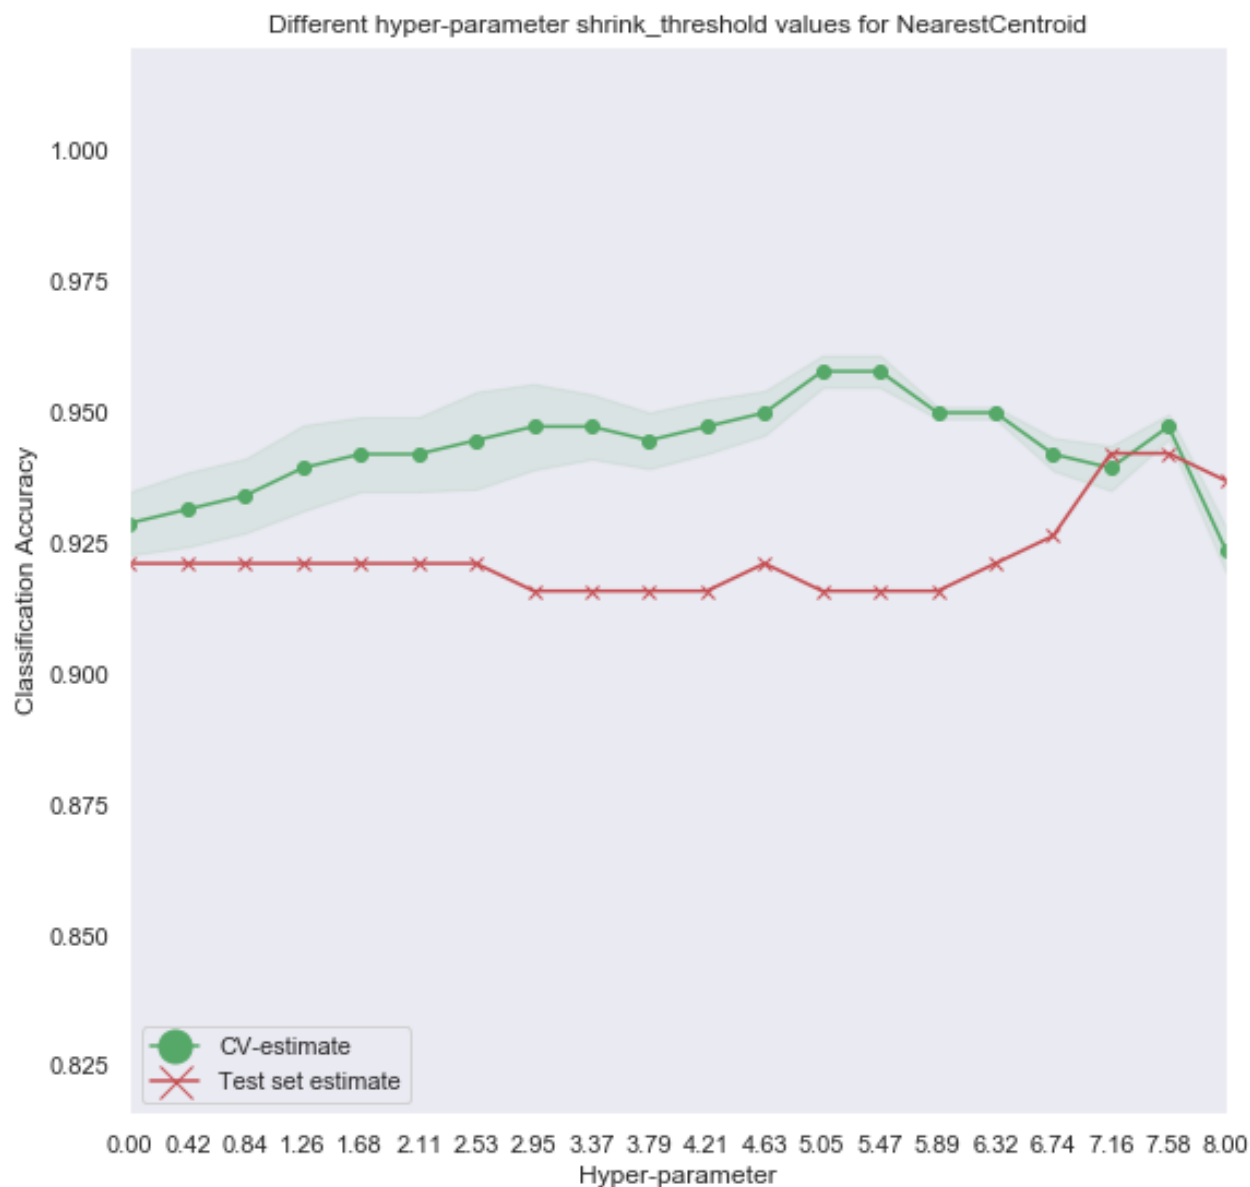
Output

1.2.1 Hyper-parameter guessing plots

Script 1.2.2 (python)

```
1 gp_disable_plots = False
2 gp_verbose = False
3 gp_dim_red = False
4 gp_retained_variance = 99
5 gp_qda_hyper = True
6 gp_nsc_hyper = True
7 gp_iterations = 1
8 gp_show_progress = False
9 gp_best_hyper_method = "max_in_cv"
10 ge_features_reduction = 0
11
12 learn_dataset(data_file = './data/wdbc.csv', header = None, random_state=0,
13               label_pos=1, label_value="M", features_ini = 2, features_fin = None,
14               best_param_value_method = gp_best_hyper_method,
15               with_dim_red = gp_dim_red, retained_variance = gp_retained_variance)
```



Output

Parameters

	Parameters
QDA hyper-parameter reg_param	True
NSC Hyper-parameter shrink_threshold	True
Dimensionality Reduction	False
Retained variance	99
Number of iterations	1
Test set size	33.3333

Best hyperparameter select method max_in_cv

Output

Accuracy

	NBC	LDA	QDA	NSC
Train	(0.95, 0.0)	(0.96, 0.0)	(0.98, 0.0)	(0.94, 0.0)
Test	(0.9, 0.0)	(0.97, 0.0)	(0.96, 0.0)	(0.92, 0.0)
Overfit degree	49.9	11	18.8	30.5

Output

TPR

	NBC	LDA	QDA	NSC
Train	(0.9, 0.0)	(0.91, 0.0)	(0.94, 0.0)	(0.91, 0.0)
Test	(0.85, 0.0)	(0.93, 0.0)	(0.91, 0.0)	(0.87, 0.0)
Overfit degree	55.2	18.4	34.6	46.3

Output

TNR

	NBC	LDA	QDA	NSC
Train	(0.97, 0.0)	(1.0, 0.0)	(1.0, 0.0)	(0.97, 0.0)
Test	(0.93, 0.0)	(1.0, 0.0)	(0.98, 0.0)	(0.94, 0.0)
Overfit degree	49.5	4.3	12.2	24.2

Output

Execution data

	Execution Data
Number of samples	569
Number of features	30
Features reduction	0
Params GaussianNB	2
Params LinearDiscriminantAnalysis	6
Params QuadraticDiscriminantAnalysis	5
Params NearestCentroid	2

Output

1.3 Prostate cancer

Script 1.3.1 (python)

```
1 # Prostate Cancer
2 gp_disable_plots = True
3 gp_verbose = False
4 gp_dim_red = False
5 gp_retained_variance = 99
6 gp_qda_hyper = True
7 gp_nsc_hyper = True
8 gp_iterations = 20
9 gp_best_hyper_method = "max_in_cv"
10 ge_features_reduction = 0
11 gp_skfold_splits = 5
12
13 learn_dataset(data_file = './data/prostate.csv', header = 0, random_state = 1,
14               label_pos = -1, label_value = 1, features_ini = 0, features_fin = -1,
15               best_param_value_method = gp_best_hyper_method,
16               with_dim_red = gp_dim_red, retained_variance = gp_retained_variance)
17
18 gp_best_hyper_method = "max_in_cv"
19 gp_dim_red = True
20 gp_retained_variance = 99
21 ge_features_reduction = 0
22 learn_dataset(data_file = './data/prostate.csv', header = 0, random_state = 1,
23               label_pos = -1, label_value = 1, features_ini = 0, features_fin = -1,
24               best_param_value_method = gp_best_hyper_method,
25               with_dim_red = gp_dim_red, retained_variance = gp_retained_variance)
26
27 gp_best_hyper_method = "max_in_cv"
28 gp_dim_red = True
29 gp_retained_variance = 80
30 ge_features_reduction = 0
31 learn_dataset(data_file = './data/prostate.csv', header = 0, random_state = 1,
32               label_pos = -1, label_value = 1, features_ini = 0, features_fin = -1,
33               best_param_value_method = gp_best_hyper_method,
34               with_dim_red = gp_dim_red, retained_variance = gp_retained_variance)
```

Output

Parameters

	Parameters
QDA hyper-parameter reg_param	True
NSC Hyper-parameter shrink_threshold	True

Dimensionality Reduction	False
Retained variance	99
Number of iterations	20
Test set size	33.3333
Best hyperparameter select method	max_in_cv

Output

Accuracy

	NBC	LDA	QDA	NSC
Train	(0.72, 0.06)	(0.86, 0.02)	(0.14, 0.25)	(0.91, 0.02)
Test	(0.64, 0.12)	(0.85, 0.04)	(0.64, 0.1)	(0.89, 0.05)
Overfit degree	106.3	17.9	3557.9	26.5

Output

TPR

	NBC	LDA	QDA	NSC
Train	(0.68, 0.07)	(0.83, 0.03)	(0.12, 0.24)	(0.88, 0.03)
Test	(0.66, 0.12)	(0.8, 0.07)	(0.62, 0.18)	(0.85, 0.07)
Overfit degree	29.6	36.8	4243.1	26.7

Output

TNR

	NBC	LDA	QDA	NSC
Train	(0.76, 0.08)	(0.89, 0.03)	(0.16, 0.27)	(0.95, 0.03)
Test	(0.62, 0.18)	(0.9, 0.08)	(0.63, 0.23)	(0.93, 0.06)
Overfit degree	178.1	16.7	2863.6	19.8

Output

Execution data

	Execution Data
Number of samples	102
Number of features	12625

Features reduction	0
Params GaussianNB	2
Params LinearDiscriminantAnalysis	6
Params QuadraticDiscriminantAnalysis	5
Params NearestCentroid	2

Output

Parameters

	Parameters
QDA hyper-parameter reg_param	True
NSC Hyper-parameter shrink_threshold	True
Dimensionality Reduction	True
Retained variance	99
Number of iterations	20
Test set size	33.3333
Best hyperparameter select method	max_in_cv

Output

Accuracy

	NBC	LDA	QDA	NSC
Train	(0.92, 0.04)	(1.0, 0.0)	(0.87, 0.15)	(0.75, 0.08)
Test	(0.74, 0.12)	(0.9, 0.04)	(0.71, 0.09)	(0.72, 0.11)
Overfit degree	200.6	100	181.1	41.3

Output

TPR

	NBC	LDA	QDA	NSC
Train	(0.92, 0.05)	(1.0, 0.0)	(0.83, 0.19)	(0.75, 0.08)
Test	(0.95, 0.05)	(0.87, 0.07)	(0.71, 0.14)	(0.73, 0.12)
Overfit degree	40.9	128.7	140.4	22.2

Output

TNR

	NBC	LDA	QDA	NSC
Train	(0.93, 0.05)	(1.0, 0.0)	(0.92, 0.13)	(0.74, 0.1)
Test	(0.48, 0.29)	(0.93, 0.07)	(0.7, 0.19)	(0.69, 0.17)
Overfit degree	483.8	66.9	235.8	62.9

Output

Execution data

	Execution Data
Number of samples	102
Number of features	12625
Features reduction	63
Params GaussianNB	2
Params LinearDiscriminantAnalysis	6
Params QuadraticDiscriminantAnalysis	5
Params NearestCentroid	2

Output

Parameters

	Parameters
QDA hyper-parameter reg_param	True
NSC Hyper-parameter shrink_threshold	True
Dimensionality Reduction	True
Retained variance	80
Number of iterations	20
Test set size	33.3333
Best hyperparameter select method	max_in_cv

Output

Accuracy

	NBC	LDA	QDA	NSC
Train	(0.89, 0.06)	(0.96, 0.02)	(0.98, 0.01)	(0.72, 0.08)
Test	(0.78, 0.08)	(0.85, 0.05)	(0.87, 0.05)	(0.67, 0.12)
Overfit degree	116.8	110.6	116.8	76.3

Output

TPR

	NBC	LDA	QDA	NSC
Train	(0.83, 0.1)	(0.93, 0.04)	(0.97, 0.02)	(0.7, 0.09)
Test	(0.66, 0.13)	(0.82, 0.07)	(0.88, 0.07)	(0.69, 0.12)
Overfit degree	200.4	122.7	91.3	11.2

Output

TNR

	NBC	LDA	QDA	NSC
Train	(0.94, 0.03)	(0.98, 0.02)	(0.99, 0.01)	(0.75, 0.08)
Test	(0.93, 0.07)	(0.9, 0.1)	(0.86, 0.11)	(0.64, 0.2)
Overfit degree	15.2	82.4	133.7	137.2

Output

Execution data

	Execution Data
Number of samples	102
Number of features	12625
Features reduction	19
Params GaussianNB	2
Params LinearDiscriminantAnalysis	6
Params QuadraticDiscriminantAnalysis	5
Params NearestCentroid	2

Output

1.3.1 Hyper-parameter guessing plots

Script 1.3.2 (python)

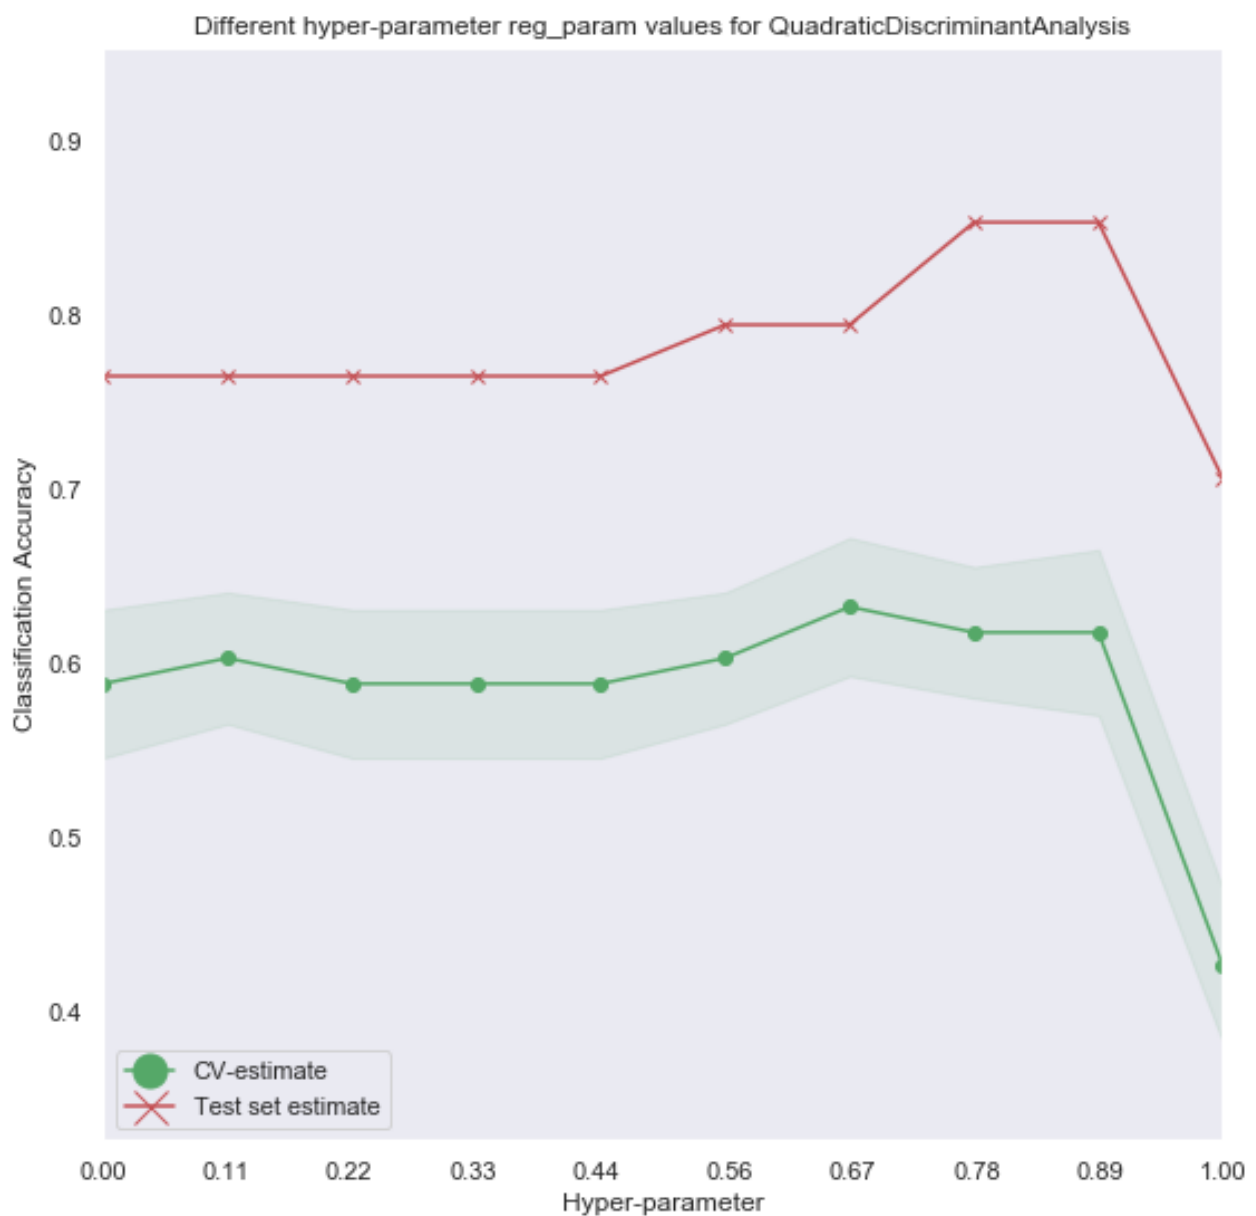
```
1 gp_disable_plots = False
2 gp_verbose = False
```

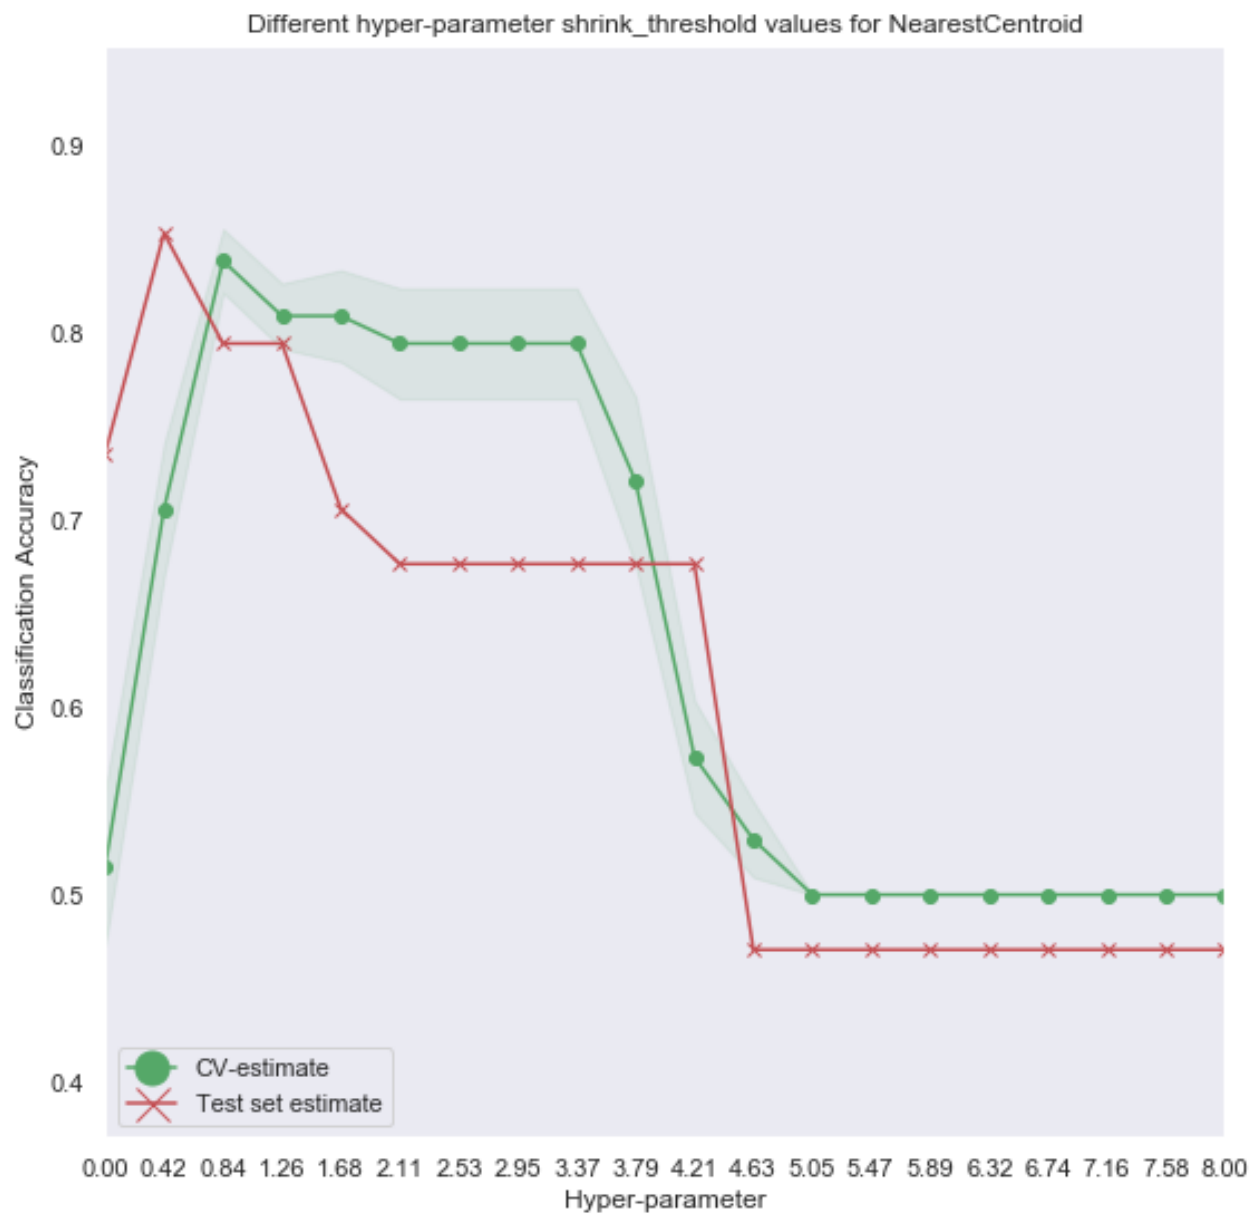


```

3 gp_dim_red = True
4 gp_retained_variance = 99
5 gp_qda_hyper = True
6 gp_nsc_hyper = True
7 gp_iterations = 1
8 gp_best_hyper_method = "max_in_cv"
9 ge_features_reduction = 0
10 gp_skfold_splits = 5
11 learn_dataset(data_file = './data/prostate.csv', header = 0, random_state = 1,
12               label_pos = -1, label_value = 1, features_ini = 0, features_fin = -1,
13               best_param_value_method = gp_best_hyper_method,
14               with_dim_red = gp_dim_red, retained_variance = gp_retained_variance)

```





Output

Parameters

	Parameters
QDA hyper-parameter reg_param	True
NSC Hyper-parameter shrink_threshold	True

Dimensionality Reduction	True
Retained variance	99
Number of iterations	1
Test set size	33.3333
Best hyperparameter select method	max_in_cv

Output

Accuracy

	NBC	LDA	QDA	NSC
Train	(0.94, 0.0)	(1.0, 0.0)	(0.93, 0.0)	(0.85, 0.0)
Test	(0.82, 0.0)	(0.91, 0.0)	(0.79, 0.0)	(0.79, 0.0)
Overfit degree	125	88.2	142.9	69

Output

TPR

	NBC	LDA	QDA	NSC
Train	(0.97, 0.0)	(1.0, 0.0)	(0.91, 0.0)	(0.85, 0.0)
Test	(0.94, 0.0)	(0.83, 0.0)	(0.94, 0.0)	(0.72, 0.0)
Overfit degree	26.9	166.7	35.8	153.3

Output

TNR

	NBC	LDA	QDA	NSC
Train	(0.91, 0.0)	(1.0, 0.0)	(0.94, 0.0)	(0.85, 0.0)
Test	(0.69, 0.0)	(1.0, 0.0)	(0.62, 0.0)	(0.88, 0.0)
Overfit degree	246	0	335.9	25.9

Output

Execution data

	Execution Data
Number of samples	102
Number of features	12625

Features reduction	63
Params GaussianNB	2
Params LinearDiscriminantAnalysis	6
Params QuadraticDiscriminantAnalysis	5
Params NearestCentroid	2

Output

1.4 Conclusions

Executions

We have two executions in breast, one without PCA and another with PCA (99 percent variance retained), and three in prostate, without PCA, with PCA 99 percent variance retention and PCA 80 percent variance. There is a parameter that indicates the level of possible overfit. The more his value, less robust to overfitting is the learning method, thus more flexible as more closed to the training data.

We have defined this *overfitting indicator* as the ratio of the difference of accuracies between train and test as numerator, and as denominator the test accuracy

The tables show the means of 20 executions for each of the four learning methods and close to them the standard deviations.

Breast

The best accuracy methods are *QDA* and *LDA*.

The reduction of dimensions is more effective in *QDA*.

According to the value of the overfit indicator, the most flexible methods are *LDA* and *QDA*, but when reducing the dimensions *NBC* goes to the first position

Prostate

We observe that *QDA* performs very poorly, as expected given the high dimensionality of this dataset (more than 12625) and the relatively low number of samples (around 100), which do not ease the accurate computation of the covariance matrices.

This affects all methods, the best ones being *NSC* and *LDA*.

The overfit of *QDA* is huge, and another order of magnitude is also relevant for *NBC*.

We also note that the accuracy of *QDA* varies greatly with the sample (high standard deviation), so we are already seeing that this method is not suitable for this dataset

Through dimensionality reduction by PCA, we have seen that we improve this results.

In fact *QDA* improves significantly and becomes reliable. In another order of magnitude *NBC* and *LDA* improve but not *NSC* because on its own way, it already reduces the complexity of the model.