

Classification Systems

Daniel Cerdán, Fernando Freire

March 22, 2019

Contents

1	Classification Systems	2
1.1	Methods	3
1.2	Breast cancer	7
1.3	Prostate cancer	9

1 Classification Systems

In this practical, you are asked to compare the prediction error of:

1. The Naive Bayes Classifier
2. LDA
3. QDA
4. Nearest Shrunken Centroids Classifier

On the Breast Cancer dataset provided in the previous notebooks, and the Prostate cancer dataset attached. The details about this last dataset are found in the reference:

Singh, D., Febbo, P., Ross, K., Jackson, D., Manola, J., Ladd, C., Tamayo, P., Renshaw, A., D'Amico, A., Richie, J., Lander, E., Loda, M., Kantoff, P., Golub, T., & Sellers, W. (2002). Gene expression correlates of clinical prostate cancer behavior. *Cancer Cell*, 1, 203–209.

This dataset is in CSV format and the last column contains the class label. The task of interest is to discriminate between normal and tumor tissue samples.

Importantly:

Use a random split of 2 / 3 of the data for training and 1 / 3 for testing each classifier. Any hyper-parameter of each method should be tuned using a grid-search guided by an inner cross-validation procedure that uses only training data. To reduce the variance of the estimates, report average error results over 20 different partitions of the data into training and testing as described above. Submit a notebook showing the code and the results obtained. Give some comments about the results and respond to these questions:

What method performs best on each dataset? What method is more flexible? What method is more robust to over-fitting?

Script 1.0.1 (python)

```
1 import warnings
2 warnings.filterwarnings("ignore")
3 %matplotlib inline
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import matplotlib.lines as mlines
8 import matplotlib as mpl
9 from matplotlib import colors
10 import seaborn as sns; sns.set()
11 import scipy.stats as stats
12 import scipy as sp
13 from scipy import linalg
14 from sklearn.naive_bayes import GaussianNB
15 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
16 from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
17 from sklearn.naive_bayes import GaussianNB
18 from sklearn.neighbors import NearestCentroid
19 from sklearn.pipeline import Pipeline
20 from sklearn.model_selection import train_test_split, RepeatedStratifiedKFold, GridSearchCV
21 from sklearn import preprocessing
22 from sklearn.metrics import accuracy_score, make_scorer, confusion_matrix
```

1.1 Methods

Script 1.1.1 (python)

```
1 def create_datasets_from_file(data_file, header, random_state, label_pos,
2                               label_value, features_ini, features_fin=None):
3     """Create training and test sets from file
4
5     Args:
6         data_file (string): Name of the data file (csv) of samples a features
7         header (string): None or position of the header (pandas read_csv parameter)
8         random_state (int): Seed for the random split (as needed for sklearn
9     → train_test_split)
10        label_pos (int): Column of the labels in data_file
11        label_value (int): Value of the label to assign internal '1' value
12        features_ini (int): First column of features in data_file
13        features_fin (int): Last column + 1 of features in data_file. If None, last
14    → column of file.
15
16    Returns:
17        (np.array): train set scaled
18        (np.array): test set scaled
19        (np.array): class labels for the train set
20        (np.array): class labels for the test set
21
22    """
23    data = pd.read_csv(data_file, header = header)
24    if features_fin == None:
25        X = data.values[ :, features_ini:].astype(np.float)
26    else:
27        X = data.values[ :, features_ini:features_fin].astype(np.float)
28    y = (data.values[ :, label_pos ] == label_value).astype(np.int)
29    # Split dataset between training and test
30    x_train, x_test, y_train, y_test = train_test_split(X, y,
31                                                         test_size=1.0/3,
32                                                         → random_state=random_state)
33
34    # Data standardization
35    scaler = preprocessing.StandardScaler().fit(x_train)
36    x_train_scaled = scaler.transform(x_train)
37    x_test_scaled = scaler.transform(x_test)
38
39    # Check standardization
40    for i in range (1, np.size(x_train_scaled,1)):
41        assert round(np.var(x_train_scaled[:,0]),3) == round(np.var(x_train_scaled[:,i]),3),\
42            "Warning: revise data standardization"
43
44    return x_train_scaled, x_test_scaled, y_train, y_test
45
46 def prediction_accuracy(x_train, x_test, y_train, y_test, method_func, method_param,
47 → param_value):
48     """Estimate parameter given training and test sets:
49
50     Args:
51         x_train (np.array): train set
```

```

46         x_test (np.array): test set
47         y_train (np.array): class labels for the train set
48         y_test (np.array): class labels for the test set
49         method_func (string) : name of the learning method
50         param (string): name of learning method parameter
51         param_value (float): value of parameter to try
52     Returns:
53         float: best parameter value to use in prediction
54
55     """
56     if method_param != "" :
57         params = {method_param : param_value}
58     else:
59         params ={}
60     method = globals()[method_func](**params)
61
62     # Training
63     method.fit(x_train, y_train)
64
65     # Prediction
66     y_pred = method.predict(x_test)
67     conf = confusion_matrix(y_test, y_pred)
68     TN = conf[0][0]
69     TP = conf[1][1]
70     FP = conf[0][1]
71     FN = conf[1][0]
72     print(conf)
73     print('Prediction accuracy is: %f' % ((TP + TN) / (TN + TP + FP + FN)))
74     print('True positive rate is: %f' % (TP / (TP + FN)))
75     print('True negative rate is: %f\n' % (TN / (TN + FP)))
76
77 def estimate_parameter(x_train, x_test, y_train, y_test, method_func, param, param_values):
78     """Estimate parameter given training and test sets:
79     Args:
80         x_train (np.array): train set
81         x_test (np.array): test set
82         y_train (np.array): class labels for the train set
83         y_test (np.array): class labels for the test set
84         method_func (string) : name of the learning method
85         param (string): name of learning method parameter
86         param_values (list of float): list of parameter values to try
87     Returns:
88         (float): best parameter value to use in prediction
89
90     """
91     # Pipeline for estimate the regularization parameter
92     pipeline = Pipeline([ ('method', globals()[method_func]()) ])
93
94     # Construct the grid the hyperparameter candidate shronk theshold
95     param_grid = { 'method__' + param : param_values }
96
97     # Evaluating

```

```

98 skfold = RepeatedStratifiedKFold(n_splits=10, n_repeats=1, random_state=0)
99 gridcv = GridSearchCV(pipeline, cv=skfold, n_jobs=1, param_grid=param_grid,\
100     scoring=make_scorer(accuracy_score))
101 result = gridcv.fit(x_train, y_train)
102
103 # Accuracies
104 accuracies = gridcv.cv_results_['mean_test_score']
105 std_accuracies = gridcv.cv_results_['std_test_score']
106
107 test_accuracies = np.ones(len(param_values))
108
109 for i in range(len(param_values)):
110     method_params = {param : param_values[ i ]}
111     method = globals()[method_func](**method_params)
112     method.fit(x_train, y_train)
113     test_accuracies[ i ] = accuracy_score(method.predict(x_test), y_test)
114
115 max_test_accuracy = max(test_accuracies)
116
117 # Obtain best_param_value as max
118 best_param_value = 0
119 best_train_accuracy = 0
120 for i in range(len(param_values)):
121     if test_accuracies[ i ] == max_test_accuracy:
122         if accuracies[i] > best_train_accuracy:
123             best_train_accuracy = accuracies[i]
124             best_param_value = param_values[i]
125
126 # Plot
127 plt.figure(figsize=(15, 10))
128 line1, = plt.plot(param_values, accuracies, 'o-', color="g")
129 line2, = plt.plot(param_values, test_accuracies, 'x-', color="r")
130 plt.fill_between(param_values, accuracies - std_accuracies / np.sqrt(10), \
131     accuracies + std_accuracies / np.sqrt(10), alpha=0.1, color="g")
132 plt.grid()
133 plt.title("Different hyper-parameter " + param + " values for " + method_func)
134 plt.xlabel('Hyper-parameter')
135 plt.xticks(np.round(np.array(param_values), 2))
136 plt.ylabel('Classification Accuracy')
137 plt.ylim((min(accuracies) - 0.1, min(1.02, max(accuracies) + 0.1)))
138
139 plt.xlim((min(param_values), max(param_values)))
140 legend_handles = [ mlines.Line2D([], [], color='g', marker='o', \
141     markersize=15, label='CV-estimate'), \
142     mlines.Line2D([], [], color='r', marker='x', \
143     markersize=15, label='Test set estimate')]
144 plt.legend(handles=legend_handles, loc = 3)
145 plt.show()
146
147 print("Best param value:", best_param_value)
148 return best_param_value
149

```

```

150 def learn_dataset(data_file, header, random_state, label_pos,
151                   label_value, features_ini, features_fin=None):
152     """Learn data sets from file, methods:
153         1. The Naive Bayes Classifier
154         2. LDA
155         3. QDA
156         4. Nearest Shrunk Centroids Classifier
157     Args:
158         data_file (string): Name of the data file (csv) of samples a features
159         header (string): None or position of the header (pandas read_csv parameter)
160         random_state (int): Seed for the random split of sets (as needed for sklearn
161     ↪ train_test_split)
162         label_pos (int): Column of the labels in data_file
163         label_value (int): Value of the label to assign internal '1' value
164         features_ini (int): First column of features in data_file
165         features_fin (int): Last column + 1 of features in data_file. If None, last
166     ↪ column of file
167     """
168     X_train_scaled, X_test_scaled, y_train, y_test = \
169         create_datasets_from_file(data_file, header, random_state,
170                                   label_pos, label_value, features_ini, features_fin=None)
171     print(X_train_scaled.shape)
172
173     # Naive Bayes accuracy
174     prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test, "GaussianNB", "", "")
175
176     # LDA accuracy
177     prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test,
178     ↪ "LinearDiscriminantAnalysis", "", "")
179
180     # QDA estimate reg parameter
181     param_values = np.linspace(0, 1, 10).tolist()
182     best_param_value = estimate_parameter(X_train_scaled, X_test_scaled, y_train, y_test, \
183     ↪ "QuadraticDiscriminantAnalysis", "reg_param", param_values)
184
185     # QDA accuracy
186     # Best parameter reg value according CV estimate
187     prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test, \
188     ↪ "QuadraticDiscriminantAnalysis", "reg_param", best_param_value)
189
190     # Centroids
191     # Best parameter shrink_threshold value according CV estimate
192     param_values = np.linspace(0, 8, 20).tolist()
193     best_param_value = estimate_parameter(X_train_scaled, X_test_scaled, y_train, y_test, \
194     ↪ "NearestCentroid", "shrink_threshold", param_values)
195
196     # Centroids accuracy
197     prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test, "NearestCentroid",
198     ↪ "shrink_threshold", best_param_value)

```

1.2 Breast cancer

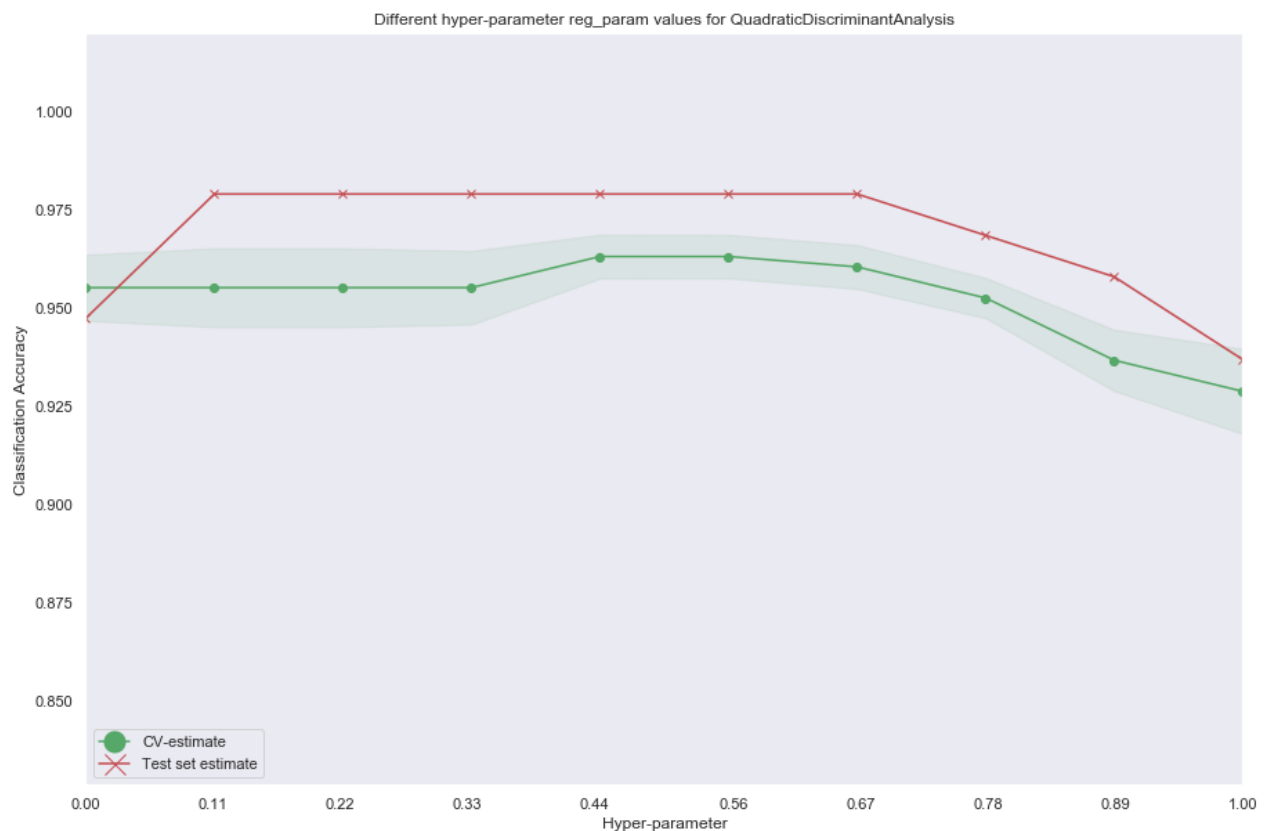
Script 1.2.1 (python)

```
1 # Breast Cancer
2 data_file = './data/wdbc.csv'
3 learn_dataset(data_file, None, 1, 1, "B", 2)
```

Output


```
(379, 30)
[[ 61  5]
 [ 7 117]]
Prediction accuracy is: 0.936842
True postive rate is: 0.943548
True negative rate is: 0.924242

[[ 60  6]
 [ 1 123]]
Prediction accuracy is: 0.963158
True postive rate is: 0.991935
True negative rate is: 0.909091
```



Output

```
Best param value: 0.4444444444444444
[[ 62  4]
 [ 0 124]]
Prediction accuracy is: 0.978947
True postive rate is: 1.000000
True negative rate is: 0.939394
```



Output

```
Best param value: 3.3684210526315788
[[ 60  6]
 [ 3 121]]
Prediction accuracy is: 0.952632
True postive rate is: 0.975806
True negative rate is: 0.909091
```


1.3 Prostate cancer


Script 1.3.1 (python)

```
1 # Prostate Cancer
2 data_file = './data/prostate.csv'
3 learn_dataset(data_file, 0, 1, -1, 1, 0, -1)
```

Output


```
(68, 12626)
[[16  0]
 [ 0 18]]
Prediction accuracy is: 1.000000
True positive rate is: 1.000000
True negative rate is: 1.000000

[[15  1]
 [ 4 14]]
Prediction accuracy is: 0.852941
True positive rate is: 0.777778
True negative rate is: 0.937500
```



Output

```
Best param value: 0.3333333333333333
[[13  3]
 [ 4 14]]
Prediction accuracy is: 0.794118
True postive rate is: 0.777778
True negative rate is: 0.812500
```



Output

```
Best param value: 2.9473684210526314
[[16  0]
 [ 0 18]]
Prediction accuracy is: 1.000000
True postive rate is: 1.000000
True negative rate is: 1.000000
```