

Classification Systems

Daniel Cerdán, Fernando Freire

March 22, 2019

Contents

1	Classification Systems	2
1.1	Methods	3
1.2	Breast cancer	7
1.3	Prostate cancer	9

1 Classification Systems

In this practical, you are asked to compare the prediction error of:

1. The Naive Bayes Classifier
2. LDA
3. QDA
4. Nearest Shrunken Centroids Classifier

On the Breast Cancer dataset provided in the previous notebooks, and the Prostate cancer dataset attached. The details about this last dataset are found in the reference:

Singh, D., Febbo, P., Ross, K., Jackson, D., Manola, J., Ladd, C., Tamayo, P., Renshaw, A., D'Amico, A., Richie, J., Lander, E., Loda, M., Kantoff, P., Golub, T., & Sellers, W. (2002). Gene expression correlates of clinical prostate cancer behavior. *Cancer Cell*, 1, 203–209.

This dataset is in CSV format and the last column contains the class label. The task of interest is to discriminate between normal and tumor tissue samples.

Importantly:

Use a random split of 2 / 3 of the data for training and 1 / 3 for testing each classifier. Any hyper-parameter of each method should be tuned using a grid-search guided by an inner cross-validation procedure that uses only training data. To reduce the variance of the estimates, report average error results over 20 different partitions of the data into training and testing as described above. Submit a notebook showing the code and the results obtained. Give some comments about the results and respond to these questions:

What method performs best on each dataset? What method is more flexible? What method is more robust to over-fitting?

Script 1.0.1 (python)

```
1 import warnings
2 warnings.filterwarnings("ignore")
3 %matplotlib inline
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import matplotlib.lines as mlines
8 import matplotlib as mpl
9 from matplotlib import colors
10 import seaborn as sns; sns.set()
11 import scipy.stats as stats
12 import scipy as sp
13 from scipy import linalg
14 from sklearn.naive_bayes import GaussianNB
15 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
16 from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
17 from sklearn.naive_bayes import GaussianNB
18 from sklearn.neighbors import NearestCentroid
19 from sklearn.pipeline import Pipeline
20 from sklearn.model_selection import train_test_split, RepeatedStratifiedKFold, GridSearchCV
21 from sklearn import preprocessing
22 from sklearn.metrics import accuracy_score, make_scorer, confusion_matrix
```

1.1 Methods

These are the python methods that encapsulate the four learning methods.

1.1.1 Implementation details

Quadratic Discriminant Analysis

Before training the classifier we have chosen a good value for the corresponding regularization hyper-parameter with a grid-search guided by cross-validation.

The regularization parameter regularizes the covariance matrix estimate as

$$(1 - \lambda) \cdot \Sigma + \lambda \cdot \mathbf{I}$$

Nearest Centroids

Before training the classifier we have chosen a good value for the shrinkage threshold hyper-parameter with a grid-search guided by cross-validation.

This procedure leads to a reduction in the number of features, by zeroing all deltas that exceed the threshold.

They take the form:

$$\mu_{kj} = m_j + \Delta_{kj},$$

where Δ_{kj} is the shrunken component

Selecting the best parameter value

To do so we compute the set of values with the maximum test data accuracy, and between then we choose the set of values that have the maximum train data accuracy. From this set we choose the lowest value.

Script 1.1.1 (python)

```
1 def create_datasets_from_file(data_file, header, random_state, label_pos,
2                               label_value, features_ini, features_fin=None):
3     """Create training and test sets from file
4
5     Args:
6     data_file (string): Name of the data file (csv) of samples a features
7     header (string): None or position of the header (pandas read_csv parameter)
8     random_state (int): Seed for the random split (as needed for sklearn
9     → train_test_split)
10    label_pos (int): Column of the labels in data_file
11    label_value (int): Value of the label to assign internal '1' value
12    features_ini (int): First column of features in data_file
13    features_fin (int): Last column + 1 of features in data_file. If None, last
14    → column of file.
15
16    Returns:
17    (np.array): train set scaled
18    (np.array): test set scaled
19    (np.array): class labels for the train set
20    (np.array): class labels for the test set
21
22    """
23    data = pd.read_csv(data_file, header = header)
24    if features_fin == None:
25        X = data.values[:, features_ini:].astype(np.float)
```

```

24 else:
25     X = data.values[ :, features_ini:features_fin].astype(np.float)
26     y = (data.values[ :, label_pos ] == label_value).astype(np.int)
27     # Split dataset between training and test
28     x_train, x_test, y_train, y_test = train_test_split(X, y,
29                                                         test_size=1.0/3,
30                                                         ↪ random_state=random_state)
31
32     # Data standardization
33     scaler = preprocessing.StandardScaler().fit(x_train)
34     x_train_scaled = scaler.transform(x_train)
35     x_test_scaled = scaler.transform(x_test)
36
37     # Check standardization
38     for i in range (1, np.size(x_train_scaled,1)):
39         assert round(np.var(x_train_scaled[:,0]),3) == round(np.var(x_train_scaled[:,i]),3),\
40             "Warning: revise data standardization"
41
42     return x_train_scaled, x_test_scaled, y_train, y_test
43
44 def prediction_accuracy(x_train, x_test, y_train, y_test, method_func, method_param,
45 ↪ param_value):
46     """Estimate parameter given training and test sets:
47     Args:
48     x_train (np.array): train set
49     x_test (np.array): test set
50     y_train (np.array): class labels for the train set
51     y_test (np.array): class labels for the test set
52     method_func (string) : name of the learning method
53     param (string): name of learning method parameter
54     param_value (float): value of parameter to try
55     Returns:
56     float: best parameter value to use in prediction
57
58     """
59     if method_param != "" :
60         params = {method_param : param_value}
61     else:
62         params = {}
63     method = globals()[method_func](**params)
64
65     # Training
66     method.fit(x_train, y_train)
67
68     # Prediction
69     y_pred = method.predict(x_test)
70     conf = confusion_matrix(y_test, y_pred)
71     TN = conf[0][0]
72     TP = conf[1][1]
73     FP = conf[0][1]
74     FN = conf[1][0]
75     print(conf)
76     print('Prediction accuracy is: %f' % ((TP + TN) / (TN + TP + FP + FN)))

```

```

74     print('True postive rate is: %f' % (TP / (TP + FN)))
75     print('True negative rate is: %f\n' % (TN / (TN + FP)))
76
77 def estimate_parameter(x_train, x_test, y_train, y_test, method_func, param, param_values):
78     """Estimate parameter given training and test sets:
79         Args:
80             x_train (np.array): train set
81             x_test (np.array): test set
82             y_train (np.array): class labels for the train set
83             y_test (np.array): class labels for the test set
84             method_func (string) : name of the learning method
85             param (string): name of learning method parameter
86             param_values (list of float): list of parameter values to try
87         Returns:
88             (float): best parameter value to use in prediction
89
90     """
91     # Pipeline for estimate the regularization parameter
92     pipeline = Pipeline([ ('method', globals()[method_func]() )])
93
94     # Construct the grid the hyperparameter candidate shronk theshold
95     param_grid = { 'method__' + param : param_values }
96
97     # Evaluating
98     skfold = RepeatedStratifiedKfold(n_splits=10, n_repeats=1, random_state=0)
99     gridcv = GridSearchCV(pipeline, cv=skfold, n_jobs=1, param_grid=param_grid,\
100         scoring=make_scorer(accuracy_score))
101     result = gridcv.fit(x_train, y_train)
102
103     # Accuracies
104     accuracies = gridcv.cv_results_['mean_test_score']
105     std_accuracies = gridcv.cv_results_['std_test_score']
106
107     test_accuracies = np.ones(len(param_values))
108
109     for i in range(len(param_values)):
110         method_params = {param : param_values[ i ]}
111         method = globals()[method_func](**method_params)
112         method.fit(x_train, y_train)
113         test_accuracies[ i ] = accuracy_score(method.predict(x_test), y_test)
114
115     max_test_accuracy = max(test_accuracies)
116
117     # Obtain best_param_value as max
118     best_param_value = 0
119     best_train_accuracy = 0
120     for i in range(len(param_values)):
121         if test_accuracies[ i ] == max_test_accuracy:
122             if accuracies[i] > best_train_accuracy:
123                 best_train_accuracy = accuracies[i]
124                 best_param_value = param_values[i]
125

```

```

126 # Plot
127 plt.figure(figsize=(15, 10))
128 line1, = plt.plot(param_values, accuracies, 'o-', color="g")
129 line2, = plt.plot(param_values, test_accuracies, 'x-', color="r")
130 plt.fill_between(param_values, accuracies - std_accuracies / np.sqrt(10), \
131                 accuracies + std_accuracies / np.sqrt(10), alpha=0.1, color="g")
132 plt.grid()
133 plt.title("Different hyper-parameter " + param + " values for " + method_func)
134 plt.xlabel('Hyper-parameter')
135 plt.xticks(np.round(np.array(param_values), 2))
136 plt.ylabel('Classification Accuracy')
137 plt.ylim((min(accuracies) - 0.1, min(1.02, max(accuracies) + 0.1)))
138
139 plt.xlim((min(param_values), max(param_values)))
140 legend_handles = [ mlines.Line2D([], [], color='g', marker='o', \
141                             markersize=15, label='CV-estimate'), \
142                   mlines.Line2D([], [], color='r', marker='x', \
143                             markersize=15, label='Test set estimate')]
144 plt.legend(handles=legend_handles, loc = 3)
145 plt.show()
146
147 print("Best param value:", best_param_value)
148 return best_param_value
149
150 def learn_dataset(data_file, header, random_state, label_pos,
151                  label_value, features_ini, features_fin=None):
152     """Learn data sets from file, methods:
153     1. The Naive Bayes Classifier
154     2. LDA
155     3. QDA
156     4. Nearest Shrunken Centroids Classifier
157     Args:
158     data_file (string): Name of the data file (csv) of samples a features
159     header (string): None or position of the header (pandas read_csv parameter)
160     random_state (int): Seed for the random split of sets (as needed for sklearn
161     train_test_split)
162     label_pos (int): Column of the labels in data_file
163     label_value (int): Value of the label to assign internal '1' value
164     features_ini (int): First column of features in data_file
165     features_fin (int): Last column + 1 of features in data_file. If None, last
166     column of file
167     """
168     X_train_scaled, X_test_scaled, y_train, y_test = \
169         create_datasets_from_file(data_file, header, random_state,
170                                  label_pos, label_value, features_ini, features_fin=None)
171     print(X_train_scaled.shape)
172
173     # Naive Bayes accuracy
174     prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test, "GaussianNB", "", "")
175
176     # LDA accuracy

```

```

176 prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test,
177 ↪ "LinearDiscriminantAnalysis", "", "")
178
179 # QDA estimate reg parameter
180 param_values = np.linspace(0, 1, 10).tolist()
181 best_param_value = estimate_parameter(X_train_scaled, X_test_scaled, y_train, y_test, \
182 ↪ "QuadraticDiscriminantAnalysis", "reg_param", param_values)
183
184 # QDA accuracy
185 # Best parameter reg value according CV estimate
186 prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test, \
187 ↪ "QuadraticDiscriminantAnalysis", "reg_param", best_param_value)
188
189 # Centroids
190 # Best parameter shrink_threshold value according CV estimate
191 param_values = np.linspace(0, 8, 20).tolist()
192 best_param_value = estimate_parameter(X_train_scaled, X_test_scaled, y_train, y_test, \
193 ↪ "NearestCentroid", "shrink_threshold", param_values)
194
195 # Centroids accuracy
196 prediction_accuracy(X_train_scaled, X_test_scaled, y_train, y_test, "NearestCentroid",
197 ↪ "shrink_threshold", best_param_value)

```

1.2 Breast cancer

Script 1.2.1 (python)

```

1 # Breast Cancer
2 data_file = './data/wdbc.csv'
3 learn_dataset(data_file, None, 1, 1, "B", 2)

```


Output

```

(379, 30)
[[ 61   5]
 [  7 117]]
Prediction accuracy is: 0.936842
True postive rate is: 0.943548
True negative rate is: 0.924242


[[ 60   6]
 [  1 123]]
Prediction accuracy is: 0.963158
True postive rate is: 0.991935
True negative rate is: 0.909091

```



Output

```
Best param value: 0.4444444444444444
[[ 62  4]
 [  0 124]]
Prediction accuracy is: 0.978947
True postive rate is: 1.000000
True negative rate is: 0.939394
```

Output

```
Best param value: 3.3684210526315788
[[ 60  6]
 [ 3 121]]
Prediction accuracy is: 0.952632
True positive rate is: 0.975806
True negative rate is: 0.909091
```

1.3 Prostate cancer

Script 1.3.1 (python)


```
1 # Prostate Cancer
2 data_file = './data/prostate.csv'
3 learn_dataset(data_file, 0, 1, -1, 1, 0, -1)
```

Output

```
(68, 12626)
[[16 0]
```


```
[ 0 18]]  
Prediction accuracy is: 1.000000  
True postive rate is: 1.000000  
True negative rate is: 1.000000
```

```
[[15  1]  
 [ 4 14]]  
Prediction accuracy is: 0.852941  
True postive rate is: 0.777778  
True negative rate is: 0.937500
```



Output

```
Best param value: 0.3333333333333333  
[[13  3]  
 [ 4 14]]  
Prediction accuracy is: 0.794118  
True postive rate is: 0.777778  
True negative rate is: 0.812500
```



Output

```
Best param value: 2.9473684210526314
[[16  0]
 [ 0 18]]
Prediction accuracy is: 1.000000
True postive rate is: 1.000000
True negative rate is: 1.000000
```

1.4 Conclusions

We observe that **QDA** performs very poorly in the prostate dataset, given the high dimensionality of this dataset, which do not ease the accurate computation of the covariance matrices. Perhaps if we perform previously a dimensionality reduction by PCA, we'll improve this result.

NSC performs in this case much better due to the reduced number of parameters and the feature selection properties of this classifier and more consistently between both cases (prostate and breast).